

## **Chapter 6: The Relational Data Model and Relational Database Constraints.**

### **Outline:**

#### 6.1 Relational Model Concepts

- Informal Definitions
- Formal Definitions
- Characteristics of Relations

#### 6.2 Keys' Types

- Super key
- Key
- Primary Key
- Alternative Key

#### 6.3 Relational Integrity Constraints

- Domain constraints
- Key constraints
- Entity integrity constraints
- Referential integrity constraints
- Semantic Integrity Constraints

#### 6.4 Update Operations and Dealing with Constraint Violations

- Insert Operation
- Delete Operation
- Update Operation

## 6.1 Relational Model Concepts

- The relational Model of Data is based on the concept of a Relation.
- A Relation is a mathematical concept based on the ideas of sets.
- The model was first introduced by Tod Codd of IBM Research in 1970 in the following paper: "A Relational Model for Large Shared Data Banks," Communications of the ACM, June 1970.

### Informal Definitions

- Informally, a relation looks like a table of values.
- A relation typically contains a set of rows.
- The data elements in each row represent certain facts that correspond to a real-world entity or relationship
  - In the formal model, rows are called tuples
- Each column has a column header that gives an indication of the meaning of the data items in that column
  - In the formal model, the column header is called an attribute

### Example of a Relation

STUDENT						
Name	Ssn	Home_phone	Address	Office_phone	Age	Gpa
Benjamin Bayer	305-61-2435	373-1616	2918 Bluebonnet Lane	NULL	19	3.21
Chung-cha Kim	381-62-1245	375-4409	125 Kirby Road	NULL	18	2.89
Dick Davidson	422-11-2320	NULL	3452 Elgin Road	749-1253	25	3.53
Rohan Panchal	489-22-1100	376-9821	265 Lark Lane	749-6492	28	3.93
Barbara Benson	533-69-1238	839-8461	7384 Fontana Lane	NULL	19	3.25

### Formal Definitions

#### 1. Schema

- The Schema (or description) of a Relation:
  - Denoted by  $R(A_1, A_2, \dots, A_n)$
  - $R$  is the name of the relation
  - The attributes of the relation are  $A_1, A_2, \dots, A_n$
- **Example:**
  - CUSTOMER (Cust-id, Cust-name, Address, Phone#)
  - CUSTOMER is the relation name
  - Defined over the four attributes: Cust-id, Cust-name, Address, Phone#
- Each attribute has a domain or a set of valid values.
- For example, the domain of Cust-id is 6 digit numbers.
- The data domain may have a data-type and/or format.

## 2. Tuple

- A tuple is an ordered set of values (enclosed in angled brackets '< ... >')
- Each value is derived from an appropriate domain.
- A *tuple* can be considered as a set of (<attribute>, <value>) pairs. Thus the following two tuples are *identical*:

```
t1 = <(Name, B. Bayer), (SSN, 305-61-2435), (HPhone, 3731616), (Address, 291 Blue Lane), (WPhone, null), (Age, 23), (GPA, 3.25)>
```

```
t2 = <(HPhone, 3731616), (WPhone, null), (Name, B. Bayer), (Age, 23), (Address, 291 Blue Lane), (SSN, 305-61-2435), (GPA, 3.25)>
```

## 3. State

- The relation state is a subset of the Cartesian product of the domains of its attributes
- Each domain contains the set of all possible values the attribute can take.
- Example: attribute Cust-name is defined over the domain of character strings of maximum length 25  
dom(Cust-name) is varchar(25)

## 4. Summary

<u>Informal Terms</u>	<u>Formal Terms</u>
Table	Relation
Column Header	Attribute
All possible Column Values	Domain
Row	Tuple
Table Definition	Schema of a Relation
Populated Table	State of the Relation

**Table 6.1:** Formal / Informal Summary

### Characteristics of Relations:

1. Ordering of tuples in a relation  $r(R)$ : The tuples are not considered to be ordered, even though they appear to be in the tabular form.
2. Ordering of attributes in a relation schema  $R$  (and of values within each tuple): We will consider the attributes in  $R(A_1, A_2, \dots, A_n)$  and the values in  $t = \langle v_1, v_2, \dots, v_n \rangle$  to be ordered.
3. Values in a tuple:
  - All values are considered atomic (indivisible).
  - Each value in a tuple must be from the domain of the attribute for that column
    - If tuple  $t = \langle v_1, v_2, \dots, v_n \rangle$  is a tuple (row) in the relation state  $r$  of  $R(A_1, A_2, \dots, A_n)$

- Then each  $v_i$  must be a value from  $\text{dom}(A_i)$
- A special null value is used to represent values that are unknown or unapplicable to certain tuples.

## 6.2 Keys' Types

- In order to understand the Relational Integrity Constraints we need to understand the Types of keys before
  - 1. Superkey of R (SK)**
    - SK is a set of attributes of R with the following condition:
      - No two tuples in any valid relation state  $r(R)$  will have the same value for SK
      - That is, for any distinct tuples  $t_1$  and  $t_2$  in  $r(R)$ ,  $t_1[SK] \neq t_2[SK]$
  - 2. Key of R**
    - A "minimal" superkey
    - That is, a key is a superkey K such that removal of any attribute from K results in a set of attributes that is not a superkey (does not possess the superkey uniqueness property)
    - Each of the keys is called **candidate key**
  - 3. Primary key (PK)**
    - Choice **one** of the candidate key
    - It is usually better to choose a primary key with a single attribute or a small number of attributes.
  - 4. Alternative key**
    - all candidate keys except primary key

### Example (1):

Consider the EMPLOYEE relation schema:

**Employee (SSN,Empid,Fname,Mname,Lname,Add,Age)**

#### 1. Superkey of R:

{SSN}, {Empid}, {SSN,Empid}, {SSN,Fname}, {SSN,Mname}, {SSN,Lname}, {SSN,Add}, {SSN,Age}, {SSN,Empid,Fname}, {SSN,Empid,Lname}, {SSN,Empid,Mname}, {SSN,Empid,Add}, {SSN,Empid,Age}, {Fname, Mname, Lname}, ...etc.

#### 2. Key of R:

{SSN}, {Empid}, {Fname,Mname,Lname}

#### 3. Primary key:

{SSN}

#### 4. Alternative key:

{Empid}, {Fname,Mname,Lname}

**Example (2):**

**Consider the following relation:**

X	Y	Z
1	A	20
3	B	21
4	A	15
1	C	30

**1. Superkey of R:**

$\{Z\}, \{X,Y\}, \{X,Z\}, \{Y,Z\}, [X,Y,Z]$

**2. Key of R:**

$\{Z\}, \{X, Y\}$

**3. Primary key:**

$\{Z\}$

**4. Alternative key:**

$\{X, Y\}$



Find superkeys, candidate keys, primary keys, alternative keys in the following CAR relation schema?

**CAR (State, Reg#, SerialNo, Model, Year)**



Find superkeys, candidate keys, primary keys, alternative keys in the following relation?

A	B	C	D	E
A1	Hh	101	D10	10
A2	Mm	200	D30	20
A3	Ff	50	D90	30
A4	Ll	60	D53	30
A1	Ff	70	D6	10

## 6.3 Relational Integrity Constraints

- **Constraints** are conditions that must hold on all valid relation states.

- There are five types of constraints in the relational model:
  1. **Domain constraints**
  2. **Key constraints**
  3. **Entity integrity constraints**
  4. **Referential integrity constraints**
  5. **Semantic Integrity Constraints**

## Domain Constraints

- Every value in a tuple must be from the domain of its attribute (or it could be null, if allowed for that attribute). A data type or format is also specified for each domain.
- Special case of this constraint **Constraints on Null** (not allowed null values for attribute)

## Key Constraints

- For any two distinct tuples  $t_1$  and  $t_2$  in a relation state  $r$  of  $R$  we have  $t_1[SK] \neq t_2[SK]$  (the subset of attributes  $SK$  is called a *superkey* of the relation schema  $R$ ).
- A superkey  $SK$  specifies a uniqueness constraint.

## Entity Integrity

- The primary key attributes PK of each relation\_schema  $R$  in  $S$  cannot have null values in any tuple of  $r(R)$ .  $t[PK] \neq \text{null}$  for any tuple  $t$  in  $r(R)$ 
  - This is because primary key values are used to identify the individual tuples.
  - If PK has several attributes(composite), null is not allowed in any of these attributes

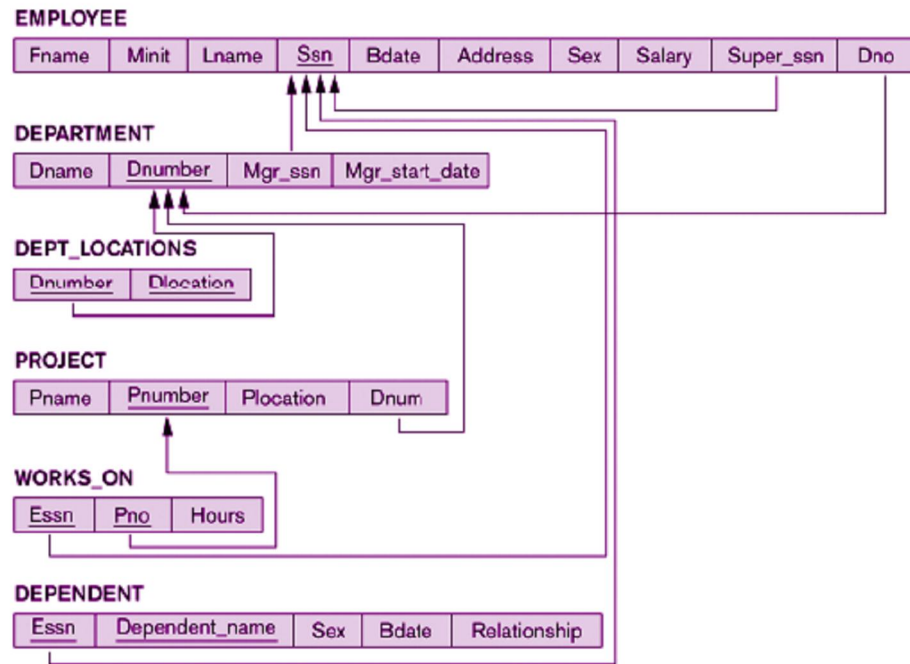
## Referential Integrity

- A constraint involving *two relations* (the previous constraints involve a *single* relation). Used to specify a *relationship* among tuples in two relations: the **referencing relation** and the **referenced relation**.
- Referential Integrity is used to maintain the consistency among tuples of the two relations.
- The referential integrity constraint states that a tuple in one relation that refers to another relation must refer to an existing tuple in that relation.
- Tuples in the *referencing relation*  $R_1$  have attributes FK (called **foreign key** attributes) that reference the primary key attributes PK of the *referenced relation*  $R_2$ . A tuple  $t_1$  in  $R_1$  is said to **reference** a tuple  $t_2$  in  $R_2$  if

$$t_1[FK] = t_2[PK]$$

- A referential integrity constraint can be displayed in a relational database schema as a directed arc from  $R_1$  FK to  $R_2$  PK.

Referential integrity constraints displayed on the COMPANY relational database schema.



- The value in the foreign key column (or columns) FK of the referencing relation R1 can be either:
  1. A value of an existing primary key value of a corresponding primary key PK in the referenced relation R2, or
  2. A null (if it is not part of its own primary key)

### Semantic Integrity Constraints

- Based on application semantics and cannot be expressed by the model.
- Example1: the max. no. of hours per employee for all projects he or she works on is 56 hrs per week.
- Example2: the salary of an employee should not exceed the salary of the employee's supervisor.
- SQL-99 allows triggers and ASSERTIONS to express for some of these

## 6.4 Update Operations and Dealing with Constraint Violations

- Operations of the relational model:
  1. Updates:
    - *Insert* - to insert a new tuple or tuples in a relation,
    - *Delete* - to delete tuples from a relation, and
    - *Update* (or *Modify*) - to change the values of some attributes in existing tuples.
  2. Retrievals - these operations do not change the current state of the databases.

- Integrity constraints should not be violated by the update operations: INSERT, DELETE, MODIFY
- In **case of integrity violation**, several actions can be taken, such as:
  - Cancel the operation that causes the violation (RESTRICT or REJECT option)
  - Trigger additional updates so the violation is corrected (CASCADE option, SET NULL option)

## Insert Operation

- ***INSERT may violate any of the constraints:***
  - Domain constraint: if one of the attribute values provided for the new tuple is not of the specified attribute domain.
  - Key constraint: if the value of a key attribute in the new tuple already exists in another tuple in the relation
  - Entity integrity: if the primary key value is null in the new tuple
  - Referential integrity: if a foreign key value in the new tuple references a primary key value that does not exist in the referenced relation

## Delete Operation

- DELETE may violate only referential integrity: If the primary key value of the tuple being deleted is referenced from other tuples in the database
- Can be remedied by several actions: RESTRICT, CASCADE, SET NULL
  - *RESTRICT option*: reject the deletion
  - *CASCADE option*: delete the tuples that reference the tuple that is being deleted
  - *SET NULL option*: set the foreign keys of the referencing tuples to NULL
- One of the above options must be specified during database design for each foreign key constraint

## Update Operation

- It may violate any of the constraints.





Which **relational constraints** would be **violated** if the following operation were applied to the COMPANY database snapshot illustrated below?

FName	LName	<u>SSN</u>	BDate	Salary	SuperSSN	DNo
Smith	Jone	123	1/2/1979	2100JD	NULL	10
Susan	Adams	432	3/6/1982	1547JD	123	10
James	Ford	587	20/9/1980	3650JD	983	20
Sali	Miller	983	1/12/1970	890JD	123	30

### Department

<u>DNo</u>	DName
10	Headquarter
20	Accounting
30	IT

### Insert

2. Insert < 'Jasmine', 'Rani', 321, 3/9/1982, 3000, NULL, 30> into EMPLOYEE.
3. Insert < 'Jasmine', 'Rani', 923, 3/19/1980, '3000JD', NULL, 30> into EMPLOYEE.
4. Insert <123456, 'Rani', 321, 3/9/1985, '3000\$', 432, 10> into EMPLOYEE.
5. Insert <400,'Planning'> into DEPARTMENT.
6. Insert <123456, 'Rani', 123, 3/9/1985, '3500JD', 432, 10> into EMPLOYEE.
7. Insert <30,'Planning'> into DEPARTMENT.

8. Insert < 'Jasmine', 'Rani', NULL, 30/9/1977, '4000JD', 587, 10> into EMPLOYEE.
9. Insert < NULL, 'Planning'> into DEPARTMENT.
10. Insert < 'Jasmine', 'Rani', 222, 3/2/1980, 3000, 587, 50> into EMPLOYEE .
11. Insert < 'Alice', 'Rani', 222, 30/2/1976, 3000, 687, 10> into EMPLOYEE.
12. Insert < 'Jasmine', 'Rani', NULL, 3/22/1976, '1000JD', 587, 50> into EMPLOYEE.
13. Insert < 'Sali', 'Rani', 'NULL', 3/2/1969, '6000JD', 587, 40> into EMPLOYEE.
14. Insert < 'Jasmine', 'Rani', 123, 3/2/76, '3200JD', 587, 60> into EMPLOYEE.

### **Delete**

1. Delete the DEPARTMENT tuple with DNo =20
2. Delete the EMPLOYEE tuple with FName ='Smith'
3. Delete the EMPLOYEE tuple with SSN=432

### **Modify**

1. Update the SALARY of the EMPLOYEE tuple with SSN=432 to 2800.

2. Update the DNO of the EMPLOYEE tuple with SSN=432 to 70.
3. Update the DNO of the DEPARTMENT tuple with DName ='IT' to 40.
4. Update the DNO of the DEPARTMENT tuple with DName ='IT' to 10.
5. Update the SSN of the EMPLOYEE tuples with FName ='Smith' to NULL.
6. Update the SSN of the EMPLOYEE tuples to 566.